

- Consider a file storage application provided by a client/server system
 - Architecture
 - Single server, cluster, remote replication
 - File location
 - Well known location(s)
 - Scalability
 - Storage, bandwidth
 - Availability
 - What proportion of requests are satisfied
 - Manageability / control
 - How well can we control the service
 - (Total) Cost
 - Purchase, operation, administration, maintenance
-

- Peer-to-Peer systems offer an alternative to the traditional client/server approach
 - “exploit resources available at the edges of the Internet”
 - commodity personal computers or workstations
 - becoming increasingly possible and desirable
 - What resources?
 - storage capacity
 - stored content
 - CPUs
 - users?
 - Examine:
 - characteristics and challenges
 - middleware
 - applications
-


- Each peer contributes resources to the system
 - the aggregated set of resources is used to provide the service
- All peers in the system have the same capabilities
 - although peers may differ over time in the roles they perform or the resources they provide
- No dependence on the availability of a centralized component or service
 - What well known P2P application violated this characteristic?
- A key issue in the design of a peer-to-peer system is the placement of migratable resources in the system (and the subsequent location of these resources)
 - balancing workload, availability, overheads (e.g. latency)

- Volatility
 - owners/managers/users of peers (hosts/nodes) in a peer-to-peer system will usually not (or cannot) provide guarantees for the availability of individual peers, leading to *churn*
- Therefore, a peer-to-peer service cannot guarantee the availability of a resource
 - although a system can be designed to make the probability of failure (i.e. unavailability) small
- This is an issue called 'harvest'
 - what proportion of the normally available resources can currently be accessed?

- 3rd Generation Peer-to-Peer systems
 - 1st Generation: e.g. Napster
 - 2nd Generation: e.g. Freenet, Kazaa, Gnutella, BitTorrent
 - examples: **Pastry**, Tapestry, CAN, Chord, Kademlia
 - Peer-to-peer middleware systems
 - place resources (files, data objects, ...) on a set of computers that are widely distributed throughout the Internet
 - route messages to these resources on behalf of clients
 - relieve client applications of this responsibility
 - Resources are identified by globally unique IDs (GUIDs)
 - a secure hash from some or all of the object's state
 - (Overlay) Routing takes place in the GUID space, not the IP address space
-

- Assign resources to nodes on behalf of applications
 - Add new resources and remove existing ones
 - Load balancing: achieve an even/fair distribution of load across peers through random placement and replication of resources
 - Subsequently, locate resources on behalf of applications
 - Optimise for local interactions between nearby clients (in the network space, **not** the GUID space)
 - Add new hosts and remove existing ones
 - Hosts in a peer-to-peer system will usually not provide guarantees about availability in the same way that centralised servers do
 - Accommodate volatile host availability (churn)
 - Redistribute resources / load
 - Other issues of less interest to us: security, anonymity, freedom
-

- A. Rowstron and P. Druschel, “Pastry: Scaleable, decentralized object location and routing for large-scale peer-to-peer systems”, 2001
 - **Required reading**
- Nodes and objects are assigned 128-bit GUIDs
 - Computed using a secure hash function
 - Nodes in close proximity w.r.t. the nodeid space are likely to be distant w.r.t. the network, geographic, administrative, ... spaces
- Pastry uses ***prefix routing*** to route a message with destination D to the node that is numerically closest to D
 - in each routing step, a message is (usually) forwarded to a node B that shares with D a nodeid prefix that is b bits (if $b=4$, then 1 hex digit) longer than that shared between A and D, for example:




- Each Pastry node maintains routing state information
 - ***Routing Table R***
 - $\log_2^b N$ rows
 - $2^b - 1$ entries in each row
 - each entry contains nodeid and associated IP address
 - each entry in row n refers to a node whose nodeid shares the same first n digits but whose $n+1^{\text{th}}$ digit has one of the $2^b - 1$ other possible values other than the $n+1^{\text{th}}$ digit in this node's nodeid
 - usually there will be many nodes that could be placed in a given entry of the routing table, but we only record one of them
 - in practice, an effort is made to ensure that the node selected for a given entry in R is **close** to the current node (according to some definable proximity metric, e.g. network latency)
 - if no suitable node for a given entry is known, then the entry is left blank
 - the maximum number of rows in the routing table is $128/b$ but in practice only (approx) $\log_2^b N$ rows will be populated
-

Pastry Routing State Information

9

- Routing table example
 - $b=2$, nodeId=23102103
 - each entry also contains the IP address of the associated node but this information is not shown



- ***Leaf Set L***

- Contains the set of nodes numerically closest to the current nodeId (i.e. w.r.t. nodeId, not the proximity metric)
- $|L| / 2$ closest nodes with smaller nodeIds and $|L| / 2$ closest nodes with larger nodeIds
- Used during message routing (more later)

- ***Neighbourhood Set M***

- Contains the nodeIds and IP addresses of the $|M|$ closest to the current node, w.r.t the chosen proximity metric (e.g. network latency)
- Not normally used for routing messages but is used to maintain/improve routing locality properties (more later)

- Typical values for both $|L|$ and $|M|$ are 2^b or 2^{b+1}
-

Pastry Routing Algorithm

11

- Routing algorithm describes the actions taken by a node A on receipt of a message D to forward the message to the next node along the route (to D)
 - Some notation
 - R_i^l is the routing table entry at column i and row l
 - L_i is the i^{th} closest nodeid in the leaf set, with –ve values indicating smaller nodeids and +ve values indicating larger nodeids
 - D is the GUID of the message destination
 - D_l is the value of the l^{th} digit in D
 - $\text{shl}(A,B)$ is the length of the prefix shared between A and B , measured in digits
-

Pastry Routing Algorithm

12

```
(1)  if ( $L_{-\lfloor |L|/2 \rfloor} \leq D \leq L_{\lfloor |L|/2 \rfloor}$ ) {  
(2)    //  $D$  is within range of our leaf set  
(3)    forward to  $L_i$ , s.th.  $|D - L_i|$  is minimal;  
(4)  } else {  
(5)    // use the routing table  
(6)    Let  $l = shl(D, A)$ ;  
(7)    if ( $R_l^{D_l} \neq null$ ) {  
(8)      forward to  $R_l^{D_l}$ ;  
(9)    }  
(10)   else {  
(11)     // rare case  
(12)     forward to  $T \in L \cup R \cup M$ , s.th.  
(13)        $shl(T, D) \geq l$ ,  
(14)        $|T - D| < |A - D|$   
(15)   }  
(16) }
```

- Joining an existing Pastry network (or creating a new one if this is the first node to join)
 - `nodeId = pastryInit(Credentials, Application)`
 - `Credentials`: Application specific – used for authentication
 - `Application`: Object used for application callbacks
 - Returns `nodeId` (GUID) for the joining node
- Routing a message to a destination
 - `route(msg, key)`
 - Routes `msg` through the Pastry network to the node with the `nodeId` numerically closest to the `key`
- Application object must provide
 - `deliver(msg, key)`: called when a `msg` is received and the current `nodeId` is closest to the `key`
 - `forward(msg, key, nextId)`: called by Pastry before a `msg` with destination `key` is routed to the next node in the route identified by `nodeId` – the current node may modify the message or stop it
 - `newLeafs(leafSet)`: informs application of a change in the `leafSet`

- Arriving nodes need to populate their routing table, neighbourhood set and leaf set
- Procedure
 - Initially, an arriving node with nodeid X contacts an existing node A (which it is told about or can find, e.g. through multicast/broadcast)
 - Node A routes a **join** message to Z – the existing node that is numerically closest to the new node X – through the Pastry network
 - Each node in the route A..Z along the Pastry route to X sends its state information to X which uses the received information to initialize its state
- Initializing neighbourhood set
 - It is assumed that A was in close proximity to X so A's neighbourhood set is used by X
- Initializing leaf set
 - Z will be numerically close to X (last step in the Pastry route) so Z's leaf set is used by X

- Initializing routing table
 - Entries in the first row of X's routing table (X_0) must be in close proximity to X but are independent of X's nodeid so A_0 is a good choice of entries for X_0
 - For row X_1 , where each entry has a prefix of length 1 in common with X, we can use the entries in B_1 – the second row of B's routing table, where B is the second step in the route from A .. Z
 - Similarly, for row X_2 , where each entry has a prefix of length 2 in common with X, we can use the entries in C_2 – the third row of C's routing table, where C is the third step in the route from A .. Z
 - We continue like this until we have populated the routing table
- Finally
 - X sends its new state to every node in R, L and N to allow them to update their own state using X's information

- A pastry node is considered failed when its immediate neighbours can no longer communicate with it
 - Requires repair of leaf sets in which the failed node appeared
 - Node that discovers failure will obtain leaf set from a node close to the failed node and replace the failed node with one from this other leaf set
 - Works unless $|L| / 2$ nodes fail simultaneously
 - To repair a routing table entry R_l^d , R_l^i is contacted ($i \neq d$) and sends its own R_l^d
 - If this fails to produce a valid replacement, R_{l+1}^i is contacted ($i \neq d$), and so on
 - To repair the neighbourhood set, request Neighbourhood set from remaining neighbours and update using new information

- Entries in the Pastry routing table are chosen to be in close proximity to a node
 - Each step in a pastry route forwards a message to a node that shares a prefix with the message that is greater than or equal to that shared with the present node
 - As a result, each step moves a message closer to the destination in the nodeId space, while travelling the least possible distance (ideally?) in the proximity space
 - Pastry does not route messages to achieve a globally minimum routing distance
 - Each routing decision is based on local information only
 - But the routes that are produced are “good”
 - Why?
-

- If a message is routed from a node A to a node B at a distance d , the message cannot subsequently be routed to a node at a distance of less than d from A
 - Follows from routing procedure and (correct?) maintenance of routing information
 - The expected distance travelled by a message during each successive routing step increases exponentially from one step to the next
 - Each entry in row l of the routing table is chosen from a set containing $N/2^{bl}$ nodes
 - So, a message tends to make larger steps towards its destination in each successive routing step
-

- Routing distance (nodeId space)
 - $\text{ceil}(\log_2^b N)$ is the expected maximum number of steps in a Pastry route
 - verified experimentally
 - Effect of maintaining only partial routing tables
 - 30%-40% longer routes than complete (global) routing tables
 - Replica routing
 - Many application use Pastry to route to the nearest among the k nodes numerically closest to a given destination GUID
 - Pastry locates the nearest node 76% of the time and one of the two nearest nodes 92% of the time
-